

# An intelligent agent with a human-in-the-loop approach to support the specification of software safety requirements: A pilot study

Daniel Pérez-Morera

*CITIC, PCI*

*University of Costa Rica*

San José, Costa Rica

daniel.perezmorera@ucr.ac.cr

**Abstract**—Security requirements specification remains a challenging activity in requirements engineering, particularly because it requires specialized knowledge that is often unavailable in development teams. Recent studies have shown that large language models (LLMs) can support the elicitation of security requirements derived from functional requirements, achieving high coverage but lower relevance compared to human-generated results. In this context, this work proposes an intelligent virtual agent with a human-in-the-loop approach to support business users and technical professionals with limited experience in the specification of software security requirements.

A pilot study was conducted with 20 graduate students enrolled in the Virtual Agents course at the University of Costa Rica. The experiment compared two conditions: a control group based on security requirements generated by the LLM without human intervention, and an experimental group where participants interacted with the intelligent virtual agent to review, refine, and validate the generated requirements. The evaluation considered coverage using recall, relevance using precision, and user satisfaction using the System Usability Scale (SUS).

*Index Terms*—

## I. INTRODUCTION

The software development life cycle frequently faces multiple challenges related to requirements specification, and security requirements are often overlooked or addressed only implicitly. Poor management of security requirements can lead to unintended functionality, multiple vulnerabilities, and inefficient use of resources. In many cases, the omission of security requirements is caused by time constraints or a lack of the necessary security knowledge.

However, given this shortage of security specialists, the software engineering community has established a baseline of knowledge through guidelines, standards, and recommendations that support the process of security requirements specification. Under this context, several studies [1]–[5] have explored the use of security patterns and templates by groups of practitioners for the specification of security requirements. More recently, [6] investigated the replacement of human participants with a large language model. The results of this latter study were promising, achieving high coverage in the identification of security requirements, outperforming most human participants in previous replications. However, the model

produced more conservative results in terms of relevance, generating security requirements with lower precision than those produced by humans.

These findings suggest that while LLMs are effective in recovering a broad set of possible security requirements, they still require human supervision to improve the contextual relevance and practical usefulness of the generated results. This limitation motivated the present research, which proposes the incorporation of a human-in-the-loop approach through the creation of an embodied intelligent virtual agent that supports the requirements elicitation process.

Unlike fully automated approaches, the proposed virtual agent allows users to review, reject, refine, and validate the generated security requirements during the elicitation process, combining the coverage capabilities of LLMs with human judgment and domain understanding. In addition to evaluating the quality of the generated requirements through coverage and relevance, this study also analyzes user satisfaction using the System Usability Scale (SUS), since practical adoption depends not only on technical performance but also on usability and user acceptance.

Specifically, the general objective of this research is:

*To analyze the process of eliciting security requirements implicitly derived from functional requirements using an intelligent virtual agent, with the purpose of evaluating the coverage, relevance, and user satisfaction achieved in the process, from the perspective of a requirements engineer.*

Given that the experiment was designed as a pilot study, only one execution was conducted with 20 graduate participants. All participants had prior knowledge in computing and virtual agents. The study compares a control condition, where security requirements are generated by the LLM without human intervention, and an experimental condition, where participants interact with the intelligent virtual agent using a human-in-the-loop approach.

Based on the general objective, the research questions of this study are as follows:

- **RQ1:** Is there any difference in the coverage of security requirements generated by the model compared to those generated through the use of the intelligent virtual agent?

- **RQ2:** Is there any difference in the relevance of the security requirements generated by the model compared to those generated through the use of the intelligent virtual agent?
- **RQ3:** Is there any difference in efficiency between humans working independently and humans supported by the intelligent virtual agent?
- **RQ4:** What is the level of user satisfaction with respect to the intelligent virtual agent?

For clarity, we define the following codes to refer to the different experiments:

- *NCSU13*: Original study conducted at North Carolina State University (NCSU) in 2013 and reported in 2014 (Riaz et al., 2014b).
- *UT14*: First replication conducted at the University of Trento (UT) in 2014.
- *NCSU14*: Second replication conducted at NCSU in 2014.
- *UCR15*: Third replication conducted at the University of Costa Rica (UCR) in 2015.
- *UCR18*: Fourth replication conducted at the University of Costa Rica (UCR) in 2018.
- *UCR2026*: Fifth replication conducted at the University of Costa Rica (UCR) in 2026 using large language models (LLMs).

## II. BACKGROUND & RELATED WORK

### A. Security Objectives

Security objectives are those security characteristics or properties that are desired in a system [7]. There are six main security objectives that define the behavior that must be guaranteed or prevented: (C) Confidentiality, the degree to which the data is disclosed only as intended. (I) Integrity, the degree to which a system or component guards against improper modification or destruction of computer programs or data. (A) Availability, the degree to which a system or component is operational and accessible when required for use. (ID) Identification & Authentication, the need to establish that a claimed identity is valid for a user, process, or device. (AY) Accountability, the degree to which actions affecting software assets can be traced to the actor responsible for the action. (PR) Privacy, the degree to which an actor can understand and control how their information is used. Security objectives of software systems are not only technical aspects (from the perspective of systems development), but also operational and administrative aspects [1], [8].

### B. Security Requirements

Security requirements are those that describe the desired security behavior in a system, or are also defined as the constraints of a system, where one of these constraints can implement one or more security objectives [1], [9]. Security requirements engineering (SRE) is the process of obtaining, specifying, and analyzing security requirements for a software system. Today, there are multiple approaches to support SRE processes, some based on standards such as ISO/IEC 5055

and others based on conceptual frameworks and requirements modeling, based on security objectives, resources, assets, and any risks associated with a system [8], [10], [11].

### C. Security Requirements Elicitation

Methods for the elicitation and documentation of security requirements include both conceptual frameworks and requirements modeling approaches. Among framework-oriented methods, SQUARE [11] provides a structured reference model for organizing and coordinating the technical activities and artifacts involved in security requirements engineering (SRE), following a maturity-driven approach. Another framework-based approach conceptualizes security requirements as constraints and employs them to construct satisfaction arguments for validating the satisfaction of security requirements [12].

Modeling-based approaches for security requirements elicitation include misuse cases [13], [14], misuse activities [15], and abuse cases [16]. These techniques capture the attacker's perspective and facilitate the identification of security requirements intended to mitigate potential attack scenarios. In contrast to these modeling-based approaches, our method focuses on the use of security requirements templates as the primary mechanism for identifying security requirements. It is important to note that different approaches to security requirements identification can be applied in a complementary manner to achieve a more comprehensive analysis.

Although current methodologies have been effective, they require expert personnel who understand the security objectives proposed for the system and the respective specification of the security requirements associated with these objectives, and lack detailed advice on how to contextualize the security requirements for a system [3], [4].

### D. Security Patterns and Templates

Security requirements can be reused across systems that share common security objectives. Parameterized templates have been proposed to support reusable security requirements [17], [18], although these approaches provide limited guidance on templates applicability. Empirical studies indicate that the use of patterns, catalogs, and other structured security knowledge generally improves the performance of non-experts, with few exceptions [19]–[21]. Building on this line of work, [1] identified six core security objectives and proposed an initial set of security requirements templates that can be mapped to existing security patterns to support requirement elicitation.

### E. Virtual Agents for Security Requirements Elicitation

Although security has been a widely studied topic, we have not identified any virtual agent specifically designed to support the elicitation and specification of security requirements. However, we found some related works in which virtual agents assist security analysis processes [22], [23]. One of these works presents a conversational agent that supports users in generating datacenter network configurations from informal and ambiguous requests [23]. The system combines a domain-adapted Rasa NLU pipeline, large language models,

and semantic role labeling to interpret user intent, clarify missing information, and automatically generate machine-executable network configurations. Although this approach addresses requirement elicitation in a security-related domain, its primary focus is network configuration rather than security requirements specification. Another related work proposes a multi-agent system for real-time security analysis of software releases [22]. This approach uses a set of specialized agents responsible for static and dynamic code analysis, anomaly detection, code verification, and architecture validation. The objective is to improve software security assessment by automating the identification of high-risk releases and supporting cybersecurity experts in decision-making. However, the system focuses on post-development security analysis rather than early-stage security requirements elicitation. Although both works employ agents and address security-related problems, neither proposal includes embodied virtual agents with a corporal representation capable of interacting with stakeholders during the requirements elicitation process. This represents an important gap, particularly in contexts where human-like interaction may improve communication, trust, and requirement completeness.

### III. METHODOLOGY

This study was designed as a pilot experiment to evaluate the impact of a human-in-the-loop approach in the elicitation of software security requirements using an intelligent virtual agent. The experiment compares the performance of a fully automated large language model (LLM) against the same model supported by human interaction through a virtual embodied agent.

#### A. Virtual Agent Implementation

The virtual agent was developed using Unity in a 3D environment, where an embodied avatar was created to simulate a natural interaction with the participant. The avatar was designed with a professional appearance similar to that commonly found in software development companies, with the intention of making the interaction feel familiar and realistic<sup>1</sup>. This design decision was influenced by embodied conversational agent principles proposed by Cassell [24], where visual embodiment and conversational interaction can improve user engagement and communication.

The agent has access to multiple subsystems required for real-time interaction. For Speech-to-Text (STT), OpenAI Whisper API was used due to its low response latency, which helps maintain a fluid conversation. For Text-to-Speech (TTS), GPT-4o mini TTS was selected because, in addition to its fast audio generation, it provides high-quality and customizable voices. The conversational support provided by the agent was also inspired by relational agent concepts proposed by Bickmore and Picard [25], aiming to create a more supportive and approachable interaction experience. The system was deployed as a web-based application to ensure accessibility for all participants through a standard browser.

<sup>1</sup>[https://github.com/daperez03/seqreq\\_agent.git](https://github.com/daperez03/seqreq_agent.git)



Fig. 1. Virtual Agent

As shown in Figure 1, the agent displays only the upper half of its body, maintaining a formal and approachable visual design. This professional appearance may also influence user behavior according to the Proteus Effect [26], encouraging participants to interact more attentively during the elicitation process.

#### B. Interaction Flow

The interaction process is illustrated in Figure 2. First, the participant provides a set of functional requirements to the virtual agent. The agent sends these requirements to the security requirements elicitation model, which generates a corresponding set of security requirements. These generated requirements are then presented to the participant, who can review, reject, refine, or accept them. This interaction enables the human-in-the-loop process, where human judgment complements the automated generation of the LLM.




Fig. 2. Security requirements elicitation process

#### C. Experimental Design

This study was designed as a pilot experiment to evaluate the impact of a human-in-the-loop approach in the elicitation of software security requirements using an intelligent virtual agent. The experiment compares the performance of a fully automated large language model (LLM) against the same model supported by human interaction through a virtual embodied agent.

The independent variable of this study is the use of human intervention during the elicitation process.

- **Control Group:** Security requirements generated by the LLM without human intervention.
- **Experimental Group:** Security requirements generated by the LLM with human interaction through the intelligent virtual agent (human-in-the-loop approach).

The comparison between these two conditions allows the study to isolate the impact of human interaction mediated by the intelligent virtual agent during the security requirements elicitation process. Both conditions rely on the same underlying LLM, ensuring that the core generation capabilities remain constant across the experiment. Therefore, any observed differences in the quality, completeness, clarity, or relevance of the generated security requirements can be attributed to the presence or absence of the human-in-the-loop interaction provided by the virtual agent.

#### D. Evaluation Metrics

For this research, *recall*, *precision*, *requirements identified per minute*, and the *System Usability Scale (SUS)* were used to evaluate the coverage, relevance, efficiency, and user satisfaction associated with the generation of security requirements. These metrics were selected to assess both the quality of the generated security requirements and the participants' perception of the intelligent virtual agent.

The results related to coverage and relevance were used to compare the performance of the fully automated LLM approach against the human-in-the-loop approach supported by the virtual agent. Additionally, the satisfaction results were used to evaluate the perceived usability and acceptance of the proposed solution.

To this end, the following metrics were calculated:

- **Coverage:** recall measure with respect to the security requirements templates defined in the security requirements oracle previously developed for the study. Recall measures the proportion of correctly identified security requirements compared to the total number of requirements expected in the oracle.
- **Relevance:** precision measure with respect to the security requirements templates defined in the security requirements oracle previously developed for the study. Precision determines the proportion of generated requirements that are relevant according to the oracle.
- **Efficiency:** measured as the number of valid security requirements templates identified per minute during the elicitation process. This metric evaluates how efficiently participants identify security requirements when working independently compared to participants supported by the intelligent virtual agent.
- **Satisfaction:** measured using the System Usability Scale (SUS) proposed by Brooke (1996), which evaluates the perceived usability of the intelligent virtual agent. This instrument measures aspects such as ease of use, usefulness, confidence, and user support during the security requirements elicitation process.

#### E. Experimental Procedure

The experiment was conducted following these steps:

##### 1) Informed Consent

Participants attended either a virtual or in-person session where the informed consent form was presented. This document explained the objectives of the study, data handling procedures, and the voluntary nature of participation. After reviewing the document, participants signed it to confirm their acceptance.

##### 2) Introduction

A general explanation was provided regarding existing security patterns and previous studies related to security requirements elicitation. After this introduction, participants received a list of functional requirements from which they had to identify and specify the corresponding security requirements.

##### 3) Laboratory Session

The working environment and the intelligent virtual agent were provided to each participant. The virtual agent accompanied the participant during the generation of security requirements. Each participant had a maximum of one hour to complete the task.

##### 4) Satisfaction Survey

Once the laboratory session was completed, participants answered a satisfaction survey based on the System Usability Scale (SUS). This survey aimed to evaluate their perception of the virtual agent. A maximum of 15 minutes was allocated for completing the survey.

##### 5) Closing and Data Storage

Finally, participants were thanked for their collaboration. The collected data were securely stored, ensuring confidentiality and exclusive use for research purposes.

To ensure methodological consistency between the research questions, variables, evaluation metrics, and experimental tasks, a methodological consistency matrix was defined. As shown in Table I, the matrix summarizes the relationship between each research question and the corresponding variables, instruments, and evaluation activities used throughout the experiment. The purpose of this matrix is to provide traceability between the objectives of the study and the procedures applied during data collection and analysis.

#### REFERENCES

- [1] M. Riaz, J. King, J. Slankas, and L. Williams, "Hidden in plain sight: Automatically identifying security requirements from natural language artifacts," in *2014 IEEE 22nd International Requirements Engineering Conference (RE)*. IEEE, 2014, pp. 183–192. [Online]. Available: <http://ieeexplore.ieee.org/document/6912260/>
- [2] M. Riaz, J. King, J. Slankas, L. Williams, F. Massacci, C. Quesada-López, and M. Jenkins, "Identifying the implied: Findings from three differentiated replications on the use of security requirements templates," vol. 22, no. 4, pp. 2127–2178, 2017. [Online]. Available: <http://link.springer.com/10.1007/s10664-016-9481-1>
- [3] M. Riaz, J. Slankas, J. King, and L. Williams, "Using templates to elicit implied security requirements from functional requirements - a controlled experiment," in *Proceedings of the 8th ACM/IEEE International Symposium on Empirical Software Engineering and Measurement*. ACM, 2014, pp. 1–10. [Online]. Available: <https://dl.acm.org/doi/10.1145/2652524.2652532>

| Research Question                                                                                                                                                          | Variable     | Instrument                                                                     | Task                                                                                                        |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------|--------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------|
| Is there any difference in the coverage of security requirements generated by the model compared to those generated through the use of the intelligent virtual agent?      | Coverage     | Recall                                                                         | Generate and compare security requirements against the oracle templates.                                    |
| Is there any difference in the relevance of the security requirements generated by the model compared to those generated through the use of the intelligent virtual agent? | Relevance    | Precision                                                                      | Evaluate the relevance of generated security requirements according to the oracle.                          |
| Is there any difference in efficiency between humans working independently and humans supported by the intelligent virtual agent?                                          | Efficiency   | Number of security requirements templates in the oracle identified per minute. | Measure the rate at which participants identify valid security requirements during the elicitation process. |
| What is the level of user satisfaction with respect to the intelligent virtual agent?                                                                                      | Satisfaction | System Usability Scale (SUS) questionnaire                                     | Participants complete the SUS survey after interacting with the virtual agent.                              |

TABLE I  
METHODOLOGICAL CONSISTENCY MATRIX

- [4] M. Riaz and L. Williams, "Security requirements patterns: understanding the science behind the art of pattern writing," in *2012 Second IEEE International Workshop on Requirements Patterns (RePa)*. IEEE, 2012, pp. 29–34. [Online]. Available: <http://ieeexplore.ieee.org/document/6359977/>
- [5] A. Martinez, M. Jenkins, and C. Quesada-Lopez, "Identifying implied security requirements from functional requirements," in *2019 14th Iberian Conference on Information Systems and Technologies (CISTI)*. IEEE, 2019, pp. 1–7. [Online]. Available: <https://ieeexplore.ieee.org/document/8760631/>
- [6] D. Pérez-Morera and C. Quesada-López, "Llm-supported identification of implicit security requirements from functional requirements: An empirical study," 2026, pp. 1–10.
- [7] D. Hybertson, P. Sommerlad, M. Schumacher, F. Buschmann, and E. Fernandez-Buglioni, *Security patterns: integrating security and systems engineering*. Wiley, 2013, OCLC: 865013622.
- [8] M. Riaz, "Inferring security requirements from natural language requirements artifacts." 2016. [Online]. Available: <http://www.lib.ncsu.edu/resolver/1840.16/11378>
- [9] C. B. Haley, J. D. Moffett, R. Laney, and B. Nuseibeh, "A framework for security requirements engineering," in *Proceedings of the 2006 international workshop on Software engineering for secure systems*. ACM, 2006, pp. 35–42. [Online]. Available: <https://dl.acm.org/doi/10.1145/1137627.1137634>
- [10] P. Salini and S. Kanmani, "Elicitation of security requirements for e-health system by applying model oriented security requirements engineering (MOSRE) framework," in *Proceedings of the Second International Conference on Computational Science, Engineering and Information Technology*. ACM, 2012, pp. 126–131. [Online]. Available: <https://dl.acm.org/doi/10.1145/2393216.2393238>
- [11] N. R. Mead and T. Stehney, "Security quality requirements engineering (SQUARE) methodology," vol. 30, no. 4, pp. 1–7, 2005. [Online]. Available: <https://dl.acm.org/doi/10.1145/1082983.1083214>
- [12] C. Haley, R. Laney, J. Moffett, and B. Nuseibeh, "Security requirements engineering: A framework for representation and analysis," vol. 34, no. 1, pp. 133–153, 2008. [Online]. Available: <http://ieeexplore.ieee.org/document/4359475/>
- [13] I. Alexander, "Misuse cases: use cases with hostile intent," vol. 20, no. 1, pp. 58–66, 2003. [Online]. Available: <http://ieeexplore.ieee.org/document/1159030/>
- [14] G. Sindre and A. Opdahl, "Eliciting security requirements by misuse cases," in *Proceedings 37th International Conference on Technology of Object-Oriented Languages and Systems. TOOLS-Pacific 2000*. IEEE Comput. Soc, 2000, pp. 120–131. [Online]. Available: <http://ieeexplore.ieee.org/document/891363/>
- [15] F. A. Braz, E. B. Fernandez, and M. VanHilst, "Eliciting security requirements through misuse activities," in *2008 19th International Conference on Database and Expert Systems Applications*. IEEE, 2008, pp. 328–333. [Online]. Available: <http://ieeexplore.ieee.org/document/4624737/>
- [16] J. McDermott and C. Fox, "Using abuse case models for security requirements analysis," in *Proceedings 15th Annual Computer Security Applications Conference (ACSAC'99)*. IEEE Comput. Soc, 1999, pp. 55–64. [Online]. Available: <http://ieeexplore.ieee.org/document/816013/>
- [17] D. Firesmith, "Specifying reusable security requirements." vol. 3, no. 1, p. 61, 2014. [Online]. Available: [http://www.jot.fm/contents/issue\\_2004\\_01/column6.html](http://www.jot.fm/contents/issue_2004_01/column6.html)
- [18] A. Toval, J. Nicolás, B. Moros, and F. García, "Requirements reuse for improving information systems security: A practitioner's approach," vol. 6, no. 4, pp. 205–219, 2002. [Online]. Available: <http://link.springer.com/10.1007/PL00010360>
- [19] C. Zhang and D. Budgen, "What do we know about the effectiveness of software design patterns?" vol. 38, no. 5, pp. 1213–1231, 2012. [Online]. Available: <http://ieeexplore.ieee.org/document/5975176/>
- [20] K. Yskout, R. Scandariato, and W. Joosen, "Do security patterns really help designers?" in *Proceedings of the 37th International Conference on Software Engineering - Volume 1*, ser. ICSE '15. IEEE Press, 2015, p. 292–302.
- [21] M. De Gramatica, K. Labunets, F. Massacci, F. Paci, and A. Tedeschi, "The role of catalogues of threats and security controls in security risk assessment: An empirical study with ATM professionals," in *Requirements Engineering: Foundation for Software Quality*, S. A. Fricker and K. Schneider, Eds. Springer International Publishing, 2015, vol. 9013, pp. 98–114, series Title: Lecture Notes in Computer Science. [Online]. Available: [https://link.springer.com/10.1007/978-3-319-16101-3\\_7](https://link.springer.com/10.1007/978-3-319-16101-3_7)
- [22] S. Fedorov, "Ai-driven multi-agent system for real-time security analysis of software releases," in *2025 2nd International Conference on Intelligent Computing and Robotics (ICICR)*, 2025, pp. 1195–1199.
- [23] S. AbdElRahim, S. Ghoniemy, and M. AbouRizka, "Llm-powered network requirement engineering conversational agent for automated 6g datacenter network configuration," in *2025 International Conference on Computer and Applications (ICCA)*, 2025, pp. 1–6.
- [24] J. Cassell, J. Sullivan, S. Prevost, and E. F. Churchill, Eds., *Embodied Conversational Agents*. The MIT Press, Mar. 2000. [Online]. Available: <https://direct.mit.edu/books/book/3227/Embodied-Conversational-Agents>
- [25] T. W. Bickmore and R. W. Picard, "Establishing and maintaining long-term human-computer relationships," *ACM Trans. Comput.-Hum. Interact.*, vol. 12, no. 2, p. 293–327, Jun. 2005. [Online]. Available: <https://doi.org/10.1145/1067860.1067867>
- [26] N. Yee and J. Bailenson, "The Proteus Effect: The Effect of Transformed Self-Representation on Behavior," *Human Communication Research*, vol. 33, no. 3, pp. 271–290, Jul. 2007. [Online]. Available: <https://academic.oup.com/hcr/article/33/3/271-290/4210718>